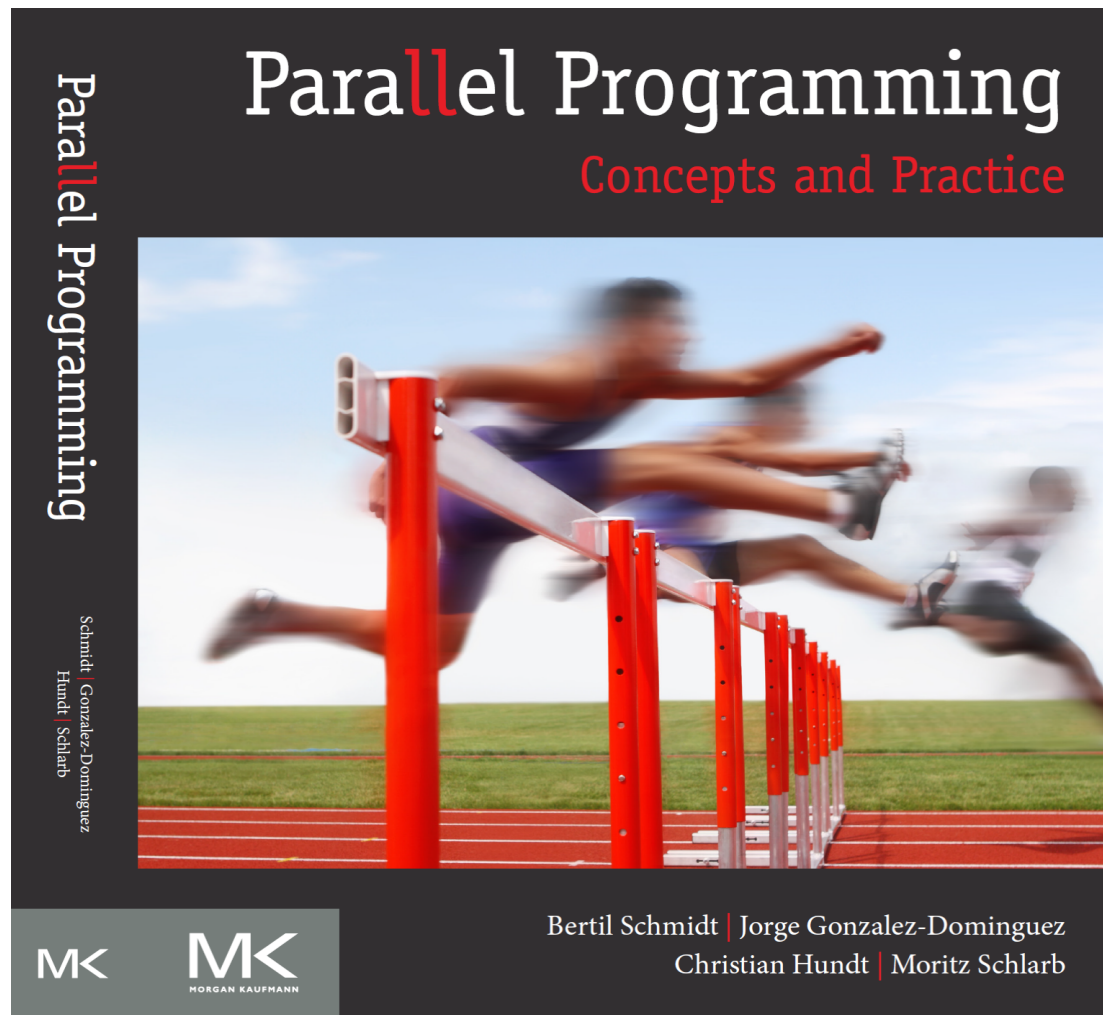
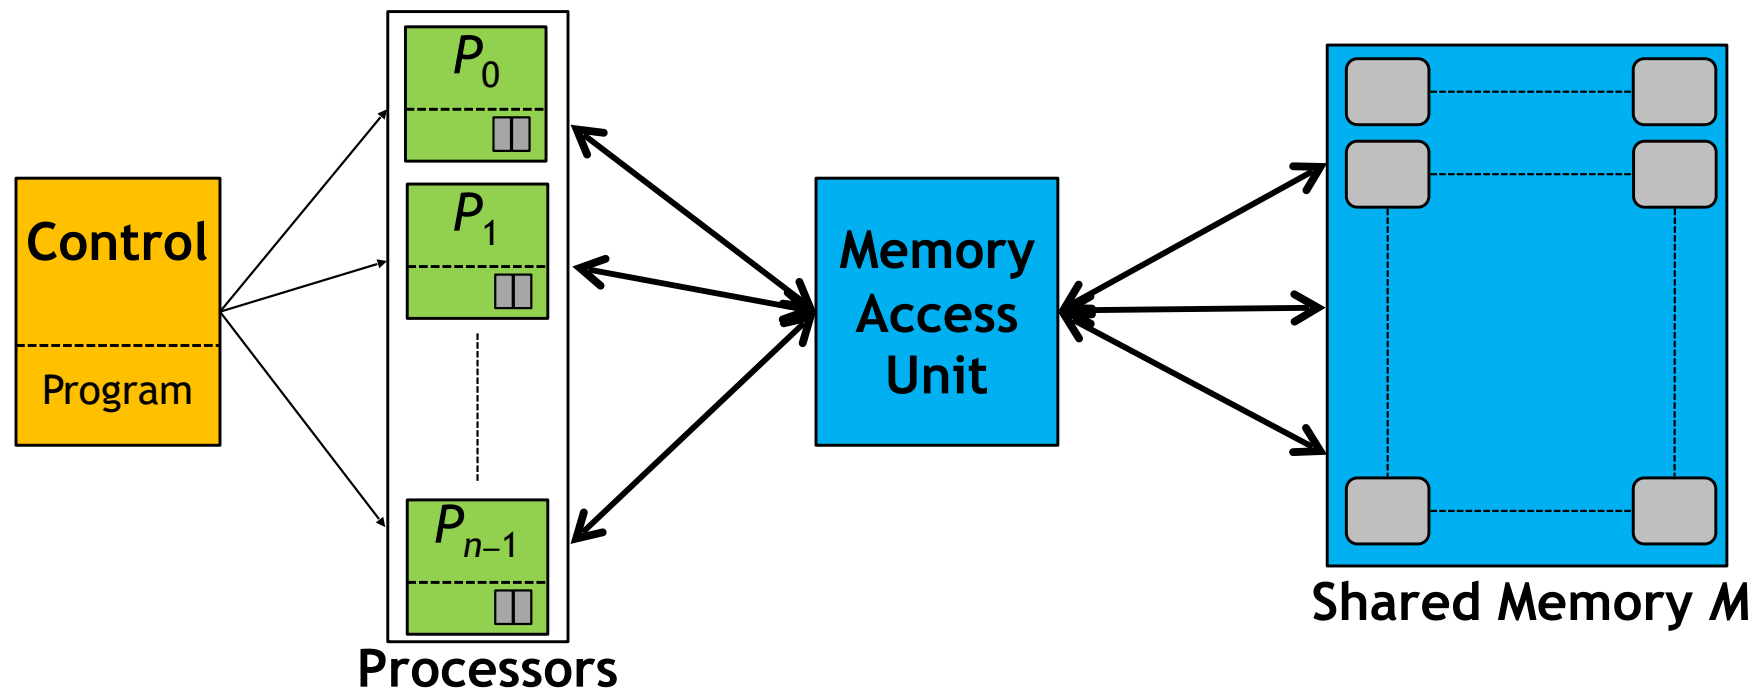


Chapter 02: Theoretical Background



Parallel Random Access Machine (PRAM)

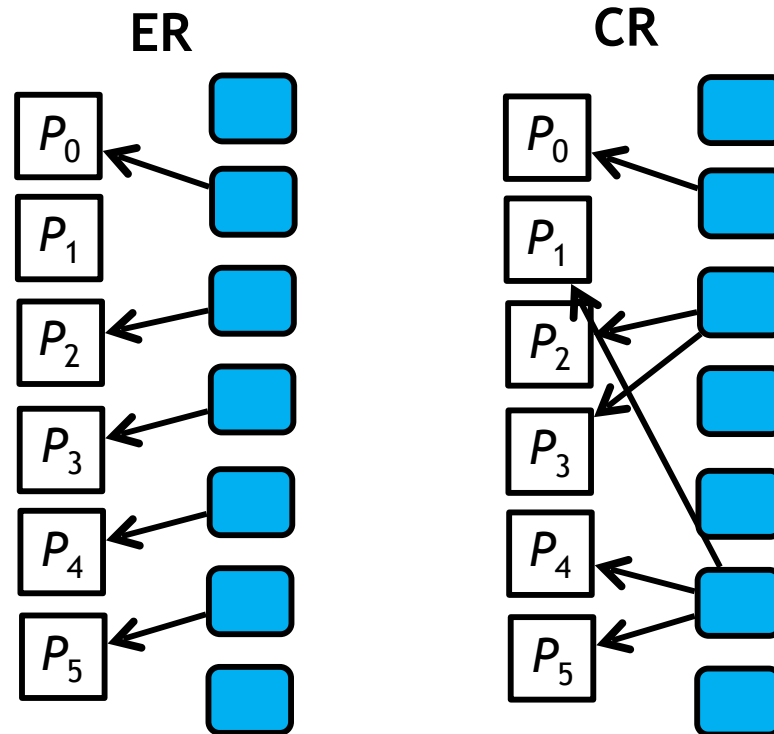


- n processors P_0, \dots, P_{n-1} are connected to a global shared memory M
- Any memory location is uniformly accessible from any processor in constant time
- Communication between processors can be implemented by reading and writing to the globally accessible shared memory.

PRAM

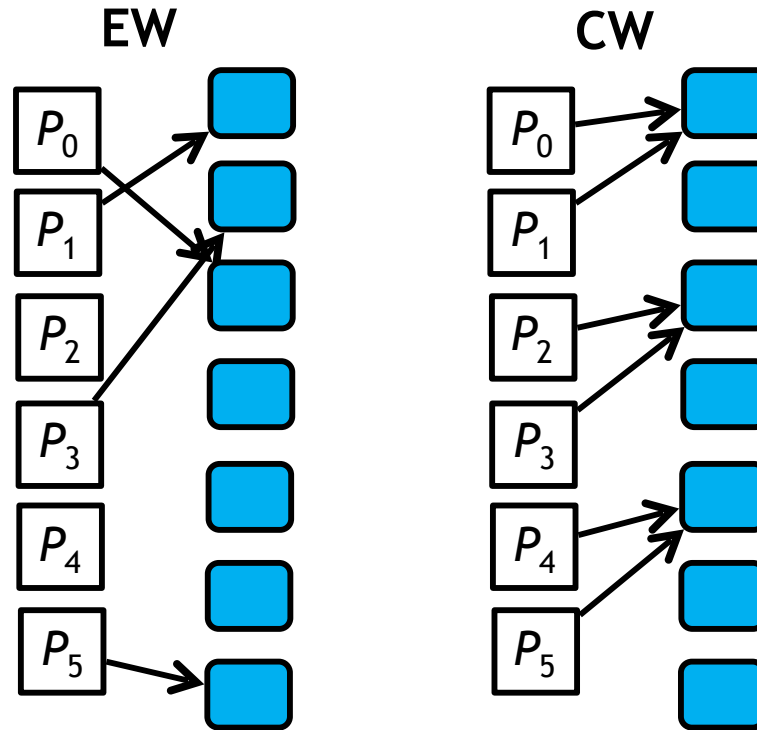
- n identical processors P_i , $i = 0, \dots, n-1$ operate in lock-step
- In every step each processor executes an instruction cycle in three phases:
 1. **Read phase:** Each processor can simultaneously read a single data item from a (distinct) shared memory cell and store it in a local register.
 2. **Compute phase:** Each processor can perform a fundamental operation on its local data and subsequently stores the result in a register.
 3. **Write phase:** Each processor can simultaneously write a data item to a shared memory cell, whereby the exclusive write PRAM variant allows writing only distinct cells while concurrent write PRAM variant also allows processors to write to the same location (race conditions).
- **Uniform complexity analysis:** Each step on the RAM takes $O(1)$ time

PRAM Variants



- **Exclusive Read Exclusive Write (EREW):** No two processors are allowed to read or write to the same shared memory cell during any cycle
- **Concurrent Read Exclusive Write (CREW):** Several processors may read data from the same shared memory cell simultaneously. Still, different processors are not allowed to write to the same shared memory cell

PRAM Variants

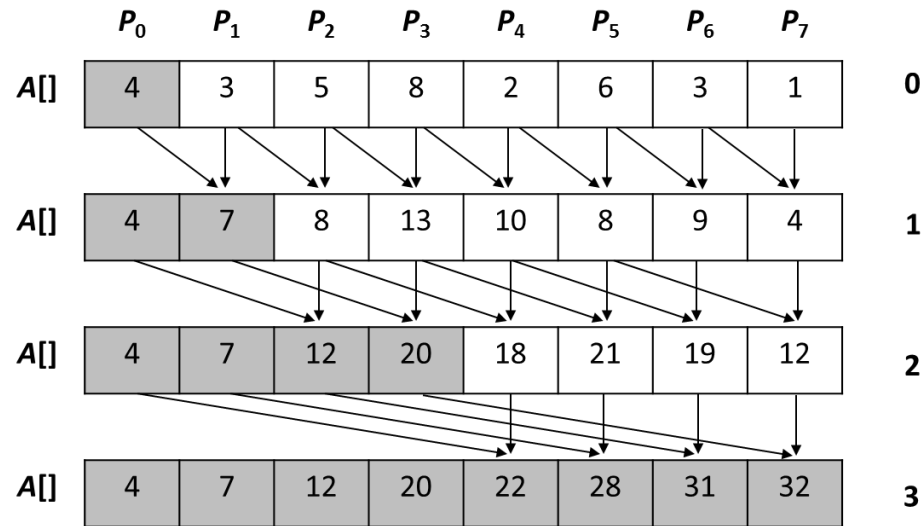


- **Concurrent Read Concurrent Write (CRCW):** Both simultaneous reads and writes to the same shared memory cell are allowed. In case of a simultaneous write (analogous to a race condition) we further specify which value will actually be stored:
 1. **Priority CW:** Processors have been assigned distinct priorities and the one with highest priority succeeds writing
 2. **Arbitrary CW:** A randomly chosen processor succeeds writing its value.
 3. **Common CW:** If the values are all equal, then this common value is written, otherwise, the memory location is unchanged.
 4. **Combining CW:** All values to be written are combined into a single value by means of an associative binary operations (e.g. sum, product, minimum, logical AND)

Parallel Prefix Computation

- Binary **associative** operation \circ on the set X ; i.e. $\circ: X \times X \rightarrow X$
 - $(x_i \circ x_j) \circ x_k = x_i \circ (x_j \circ x_k)$ for all $x_i, x_j, x_k \in X$.
 - Examples: Addition, Multiplication, Minimum, Maximum, String concatenation, boolean AND/OR
- $X = \{x_0, \dots, x_{n-1}\}$, $x_i \in X$ for all $i = 0, \dots, n-1$. We want to compute
 - $s_0 = x_0$
 - $s_1 = x_0 \circ x_1$
 - \vdots
 - $s_{n-1} = x_0 \circ x_1 \circ \dots \circ x_{n-1}$.
 - In other words, $s_0 = x_0$ and $s_i = s_{i-1} \circ x_i$ for $i = 1, \dots, n-1$
- Obtaining $S = \{s_0, \dots, s_{n-1}\}$ from $X = \{x_0, \dots, x_{n-1}\}$ is called **prefix computation**
- Example: “MINIMUM”
 - Input: {39, 21, 20, 50, 13, 18, 2, 33, 49, 39, 47, 15, 30, 47, 24, 1}
 - Output: {39, 21, 20, 20, 13, 13, 2, 2, 2, 2, 2, 2, 2, 2, 2, 1}

Parallel Prefix on a PRAM using recursive doubling on n Processors



```

for (j = 0; j < n; j++) do_in_parallel           // each processor j
    reg_j = A[j];                               // copies one value to a local register
for (i = 0; i < ceil(log(n)); i++) do           // sequential outer loop
    for (j = pow(2, i); j < n; j++) do_in_parallel // each proc. j
        reg_j += A[j - pow(2, i)];             // performs computation
        A[j] = reg_j;                          // writes result to shared memory
    }

```

- $C(n) = T(p, n) \times p = O(\log n) \times n = O(n \times \log n)$
- Thus, this algorithm is **NOT** cost-optimal

Cost-optimal Parallel Prefix on a PRAM

P0	0	1	2	3
P1	4	5	6	7
P2	8	9	10	11
P3	12	13	14	15

P0	0	1	3	6
P1	4	9	15	22
P2	8	17	27	38
P3	12	25	39	54

P0	0	1	3	6
P1	4	9	15	28
P2	8	17	27	66
P3	12	25	39	120

P0	0	1	3	6
P1	10	15	21	28
P2	36	45	55	66
P3	78	91	105	120

- Use $p = n/\log(n)$ processors
- $C(n) = T(p,n) \times p = O(\log n) \times n/\log(n) = O(n)$
- Thus, this algorithm is cost-optimal

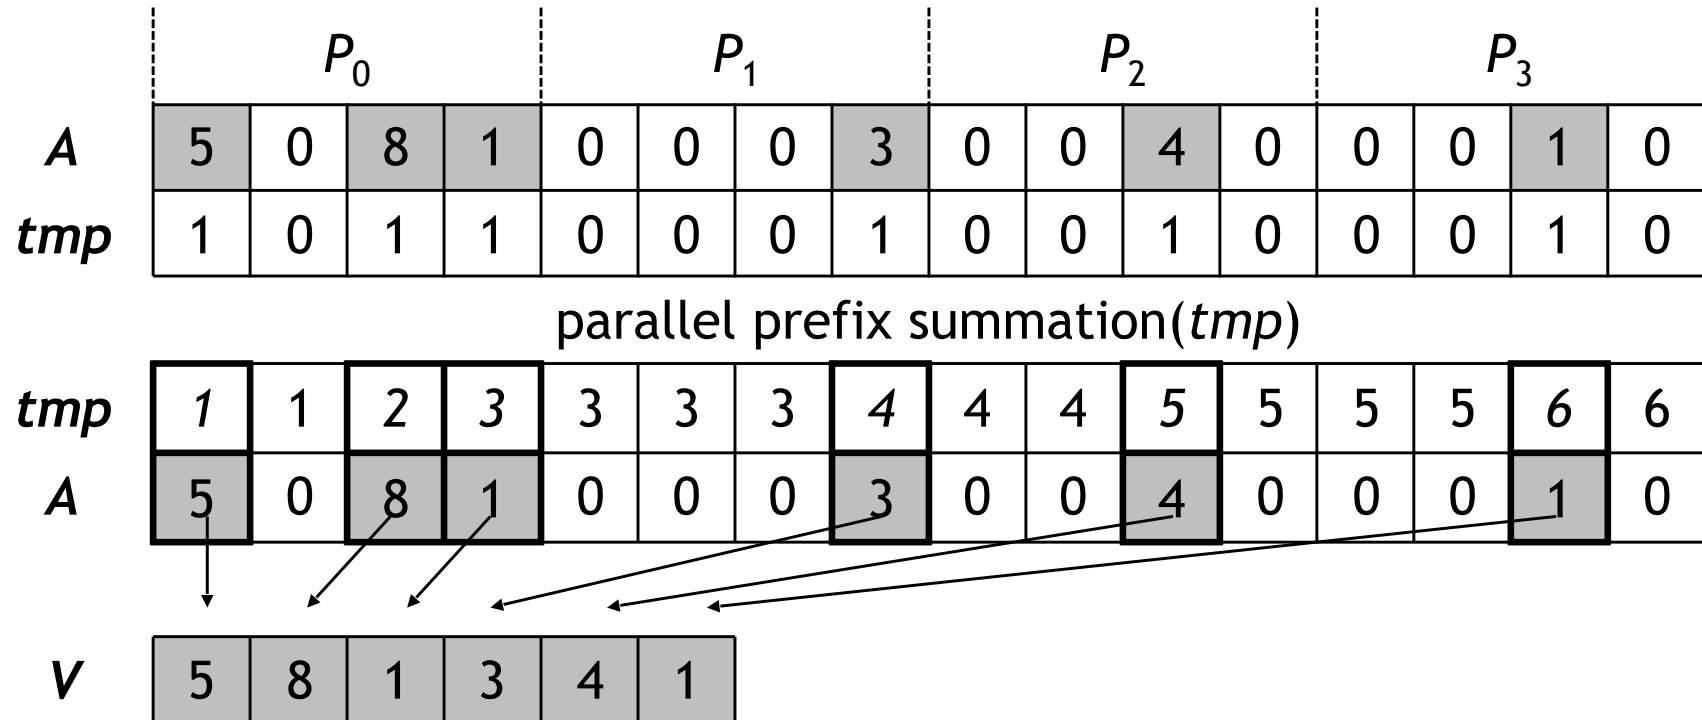
Cost-optimal Parallel Prefix on an EREW PRAM using $p = n/\log(n)$ processors

```
// Stage 1: each Processor i computes a local prefix sum
// of a subarray of size  $n/p = \log(n) = k$ 
for (i = 0; i < n / k; i++) do_in_parallel
    for (j = 1; j < k; j++) do
        A[i*k+j] += A[i*k+j-1];

// Stage 2: Prefix summation using only the rightmost value
// of each subarray ( $O(\log(n/k))$ )
for (i = 0; i < log(n / k); i++) do
    for (j = pow(2, i); j < n / k; j++) do_in_parallel
        A[j*k-1] += A[(j-pow(2, i))*k-1];

// Stage 3: each Proc i adds the value computed in Step 2 by Proc i-1 to
// each subarray element except for the last one
for (i = 1; i < n / k; i++) do_in_parallel
    for (j = 0; j < k - 1; j++) do
        A[i*k+j] += A[(i-1)*k+j];
```

Sparse Array Compaction on a PRAM

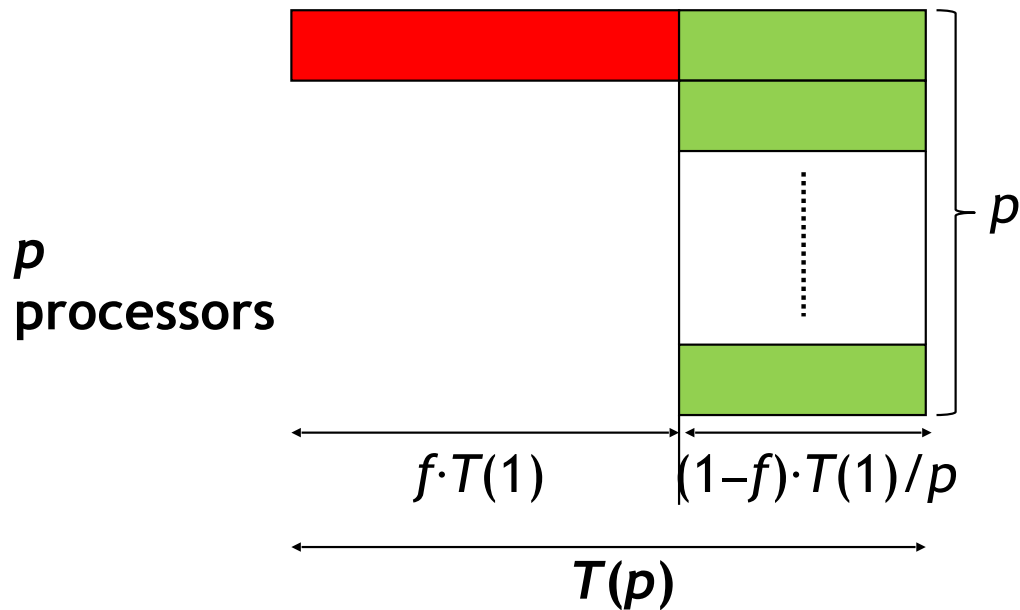
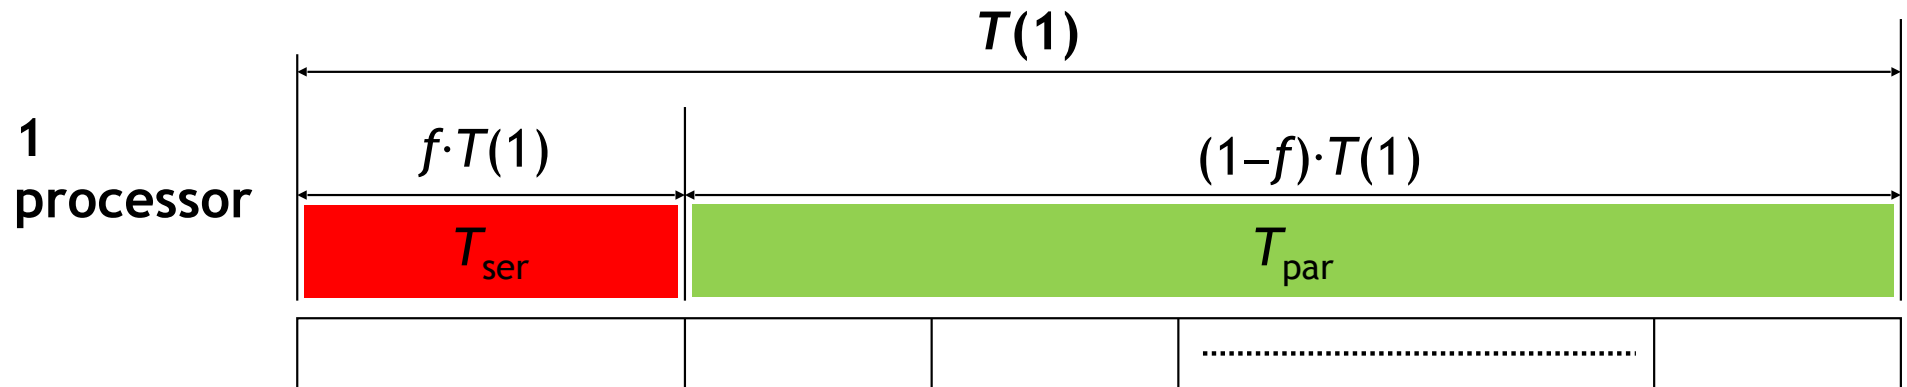


- Use $p = n/\log(n)$ processors
- $C(n) = T(p, n) \times p = O(\log n) \times n/\log(n) = O(n)$

Amdahl's Law

- A formula for estimating speedup is named Amdahl's Law
- It states that no matter how many processors are used in a parallel run, a program's speedup will be limited by its fraction of sequential code.
- That is, almost every program has a fraction of the code that doesn't lend itself to parallelism. This is the fraction of code that will have to be run with just one processor, even in a parallel run.
- Gives an **upper bound** on the speedup that can be achieved.

Amdahl's Law



$$S(p) = \frac{T(1)}{T(p)} \leq \frac{f \cdot T(1) + (1-f) \cdot T(1)}{f \cdot T(1) + \frac{(1-f) \cdot T(1)}{p}}$$

$$= \frac{f + (1-f)}{f + (1-f)/p} = \frac{1}{f + (1-f)/p}$$

Amdahl's Law

$$S(p) \leq \frac{1}{f + \frac{(1-f)}{p}}$$

- Where the term f stands for the fraction of operations done sequentially with just one processor, and the term $(1-f)$ stands for the fraction of operations that can potentially be parallelized.
- The sequential fraction of code, f , is a unit-less measure between 0 and 1.
- Amdahl's Law can be used to predict the performance of parallel programs

Amdahl's Law - Example

1. 95% of a program's execution time occurs inside a loop that can be executed in parallel. What is the maximum speedup we should expect from a parallel version of the program executing on 6 CPUs?

$$S(6) \leq \frac{1}{0.05 + \frac{(1 - 0.05)}{6}} = 4.8$$

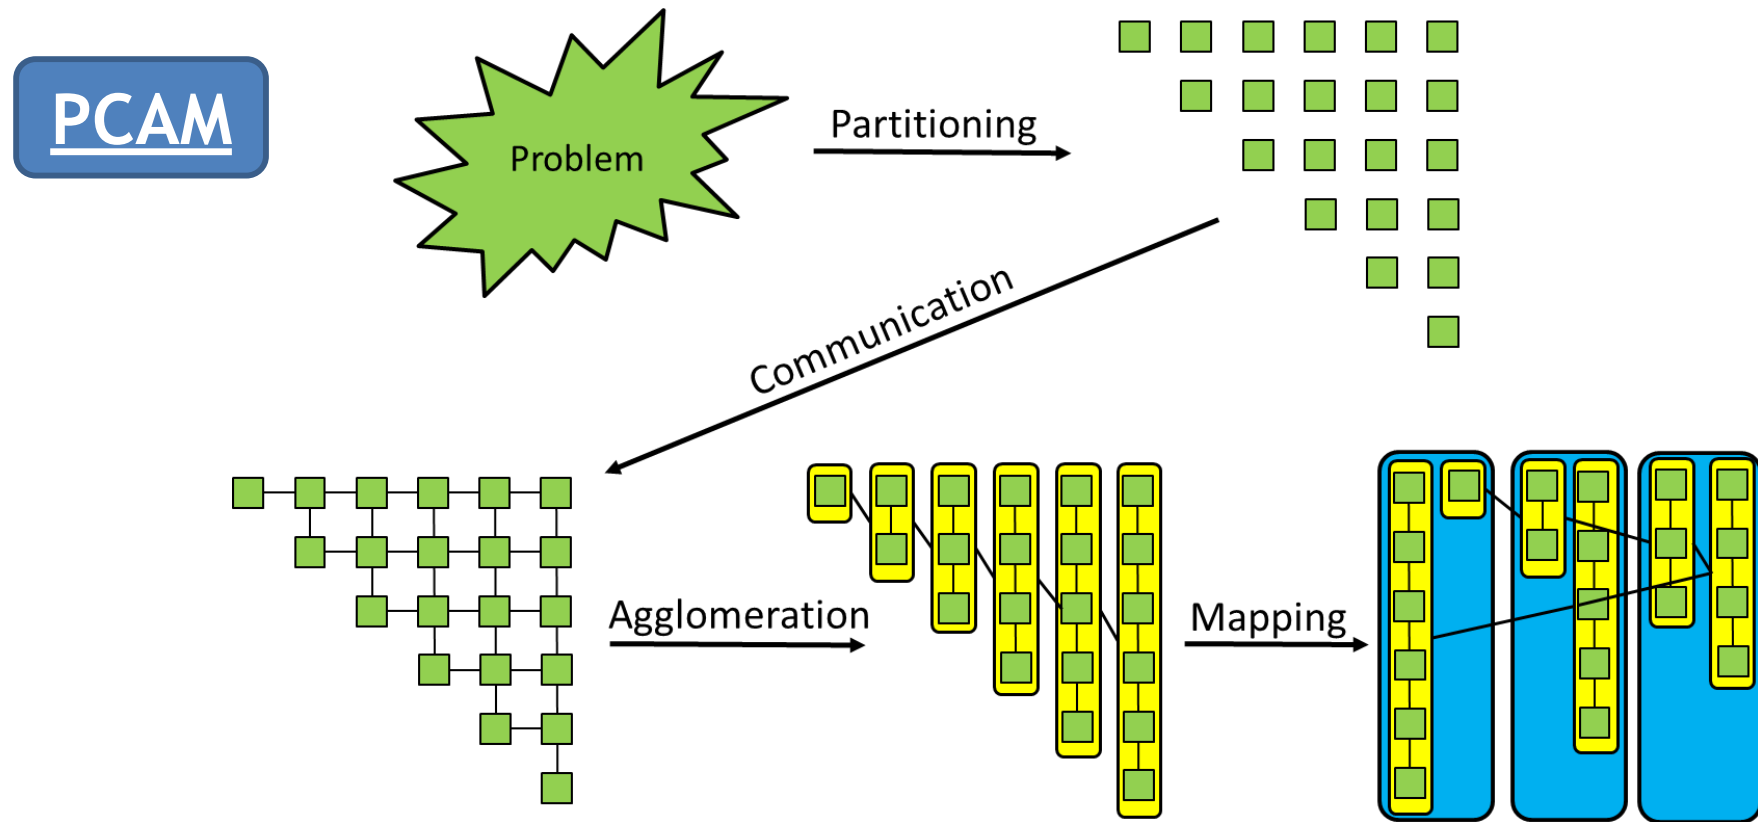
2. 10% of a program's execution time is spent within inherently sequential code. What is the limit to the speedup achievable by a parallel version of the program?

$$S(\infty) \leq \lim_{p \rightarrow \infty} \frac{1}{0.1 + \frac{(0.9)}{p}} = 10$$

Scaled Speedup

- **Limitation of Amdahl's law:** only applies in situation where the problem size is constant and the number of processors varies (\Rightarrow **strong scalability**)
- However, when using more processors we may also use larger problems sizes (\Rightarrow **weak scalability**)
- **Scaled Speedup:** incorporates such scenarios in the calculation of the achievable speedup.

Foster's Parallel Algorithm Design Method



1. **Partitioning:** decompose the problem into a large amount of small (*fine-grained*) tasks that can be executed in parallel
2. **Communication:** determine the required communication between tasks
3. **Agglomeration:** combine identified tasks into larger (*coarse-grained*) tasks in order to reduce communication by improving data locality
4. **Mapping:** assign the agglomerated to processes/threads in order to minimize communication, enable concurrency, and balance workload

Example: Jacobi Iteration

- Replace each value in the matrix by the average of its four neighbors
- Boundaries remain constant

Input Matrix: $x[i][j]$

[illegible]

xnew[i][j]

[illegible]

Jacobi Iteration

- Replaces all points of a given 2D matrix by the average of the values around it in **every iteration** step until convergence:

```
while (not converged) {  
    for (int i=1; i<rows-1; i++)  
        for (int j=1; j<cols-1; j++)  
            buff[i*cols+j] = 0.25f*(data[(i+1)*cols+j]  
                                     + data[i*cols+j-1]  
                                     + data[i*cols+j+1]  
                                     + data[(i-1)*cols+j]);  
    memcpy(data,buff,rows*cols*sizeof(float));  
}
```

- Boundary values are fixed:

```
x[0][j]  
x[n-1][j]  
x[i][0]  
x[i][n-1]
```

Example: Jacobi Iteration

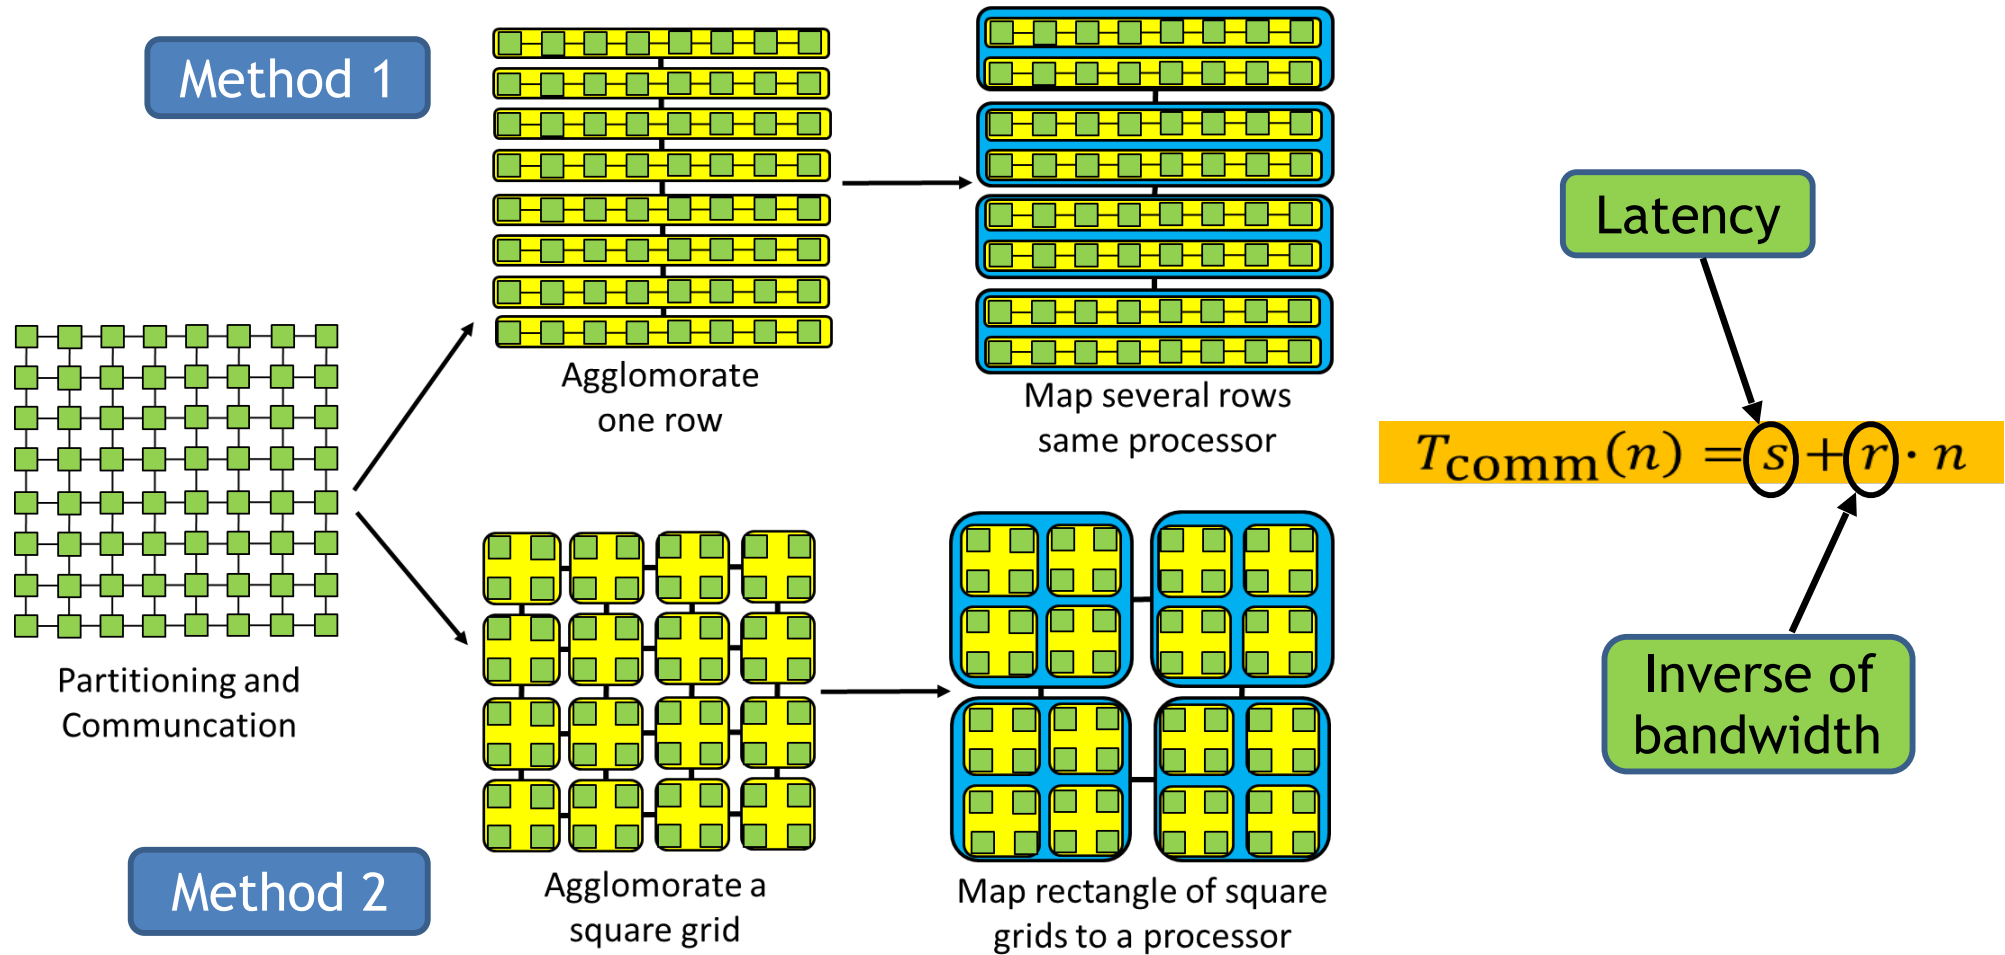
x[i][j] after 25 iterations

[illegible]

x[i][j] after 75 iterations

[illegible]

Two Schemes for Jacobi Iteration



- **Method 1:** Communication between two procs $\approx 2(s + r \cdot n)$
- **Method 2:** Communication between two procs $\approx 4 \left(s + r \left(\frac{n}{\sqrt{p}} \right) \right)$
- \Rightarrow Method 2 superior for large p since communication time decreases with p while it remains constant for Method 1.