§4.3 Scheduling

- block distribution
- cyclic "
- block-cyclic "

} false sharing issue

matrix-vector multiplication

$m \left(\boxed{A}^{n} * (x) \Rightarrow (b) \right\}^{m}$

$$b_i = \sum_{j=0}^{n-1} A_{ij} * x_j \quad , \text{for all } i \in \{0 \ldots m-1\}$$

- sequential code
- void sequential_mult(vector<double>& A, // as 1-D array like
  vector<double>& x,
  vector<double>& b, int m, int n)

```
{ for(row=0 ~ m-1)
  { accum = 0;
    for(int col=0 ~ n-1)
      accum += A[row*n + col] * x[col];
    b[row] = accum;
  }
}
```

- int main(---)

```
{ n = ...;   int n=1<<15;   // n=1 then left-shift 15 bits
  m = ...;   int m=1<<15;   // => 2^15 = 32768
  vector<double> A(m*n);
  vector<double> x(n);
  vector<double> b(m);
  // initialize A, x here
  sequential_mult(A, x, b, m, n);
  ...
}
```

vector alloc. time overhead → initialize all entries

2.7 sec vs. $1.8 * 10^{-5}$ sec.

→ less alloc. time way using dynamic array.

```
double* A = new double[m*n];
```

or,

```
vector<no_init_t<double>> A(m*n);
```

$ > g++ -O2 -std=c++11 DMV.cpp

(dense matrix vector product)

— prefix sum computation with matrix-vector mult.

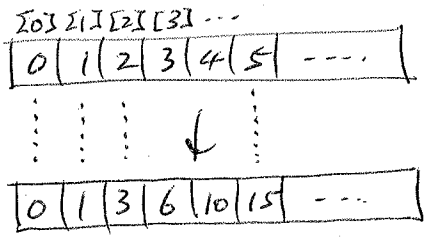

— given values — x vector

— prefix sum — b vector

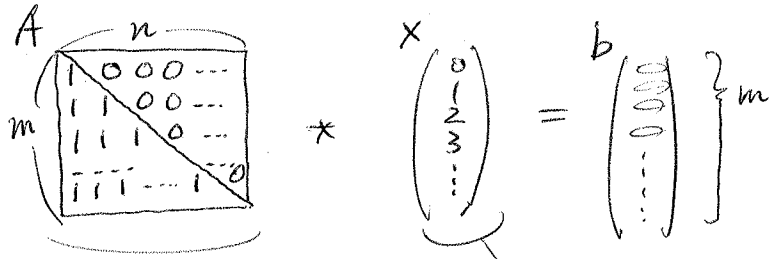

initialize A matrix as lower triangular

all 0's

all 1's

initialize x vector with any given values

```
— for (int col=0~n-1)
      X[col] = col;
```

```
— for (int row=0~m-1)
      for (int col=0~n-1)
          A[row*n+col] = row≥col ? 1 : 0; // linear storage, row-major layout.
```

— macro for get_time — inline expansion

```
#include <sys/time.h>
:
#define GET_TIME(now)\
{ struct timeval t; gettimeofday(&t, NULL);\
  now = t.tv_sec + t.tv_usec/1000000.0;\
}

:
double start_time; ---
:
GET_TIME(start_time);   // double type
:
GET_TIME(end_time);     —— cout << end_time - start_time;
```

— block distribution — version1 — main() creates threads.

DMV ex).

```
#include <cmath>
void block_parallel_mult(int id, double A[], double x[], double b[],
                         int m, int n, int num_threads)
{ int chunk = ceil((double)m / num_threads);   // for not evenly divisible case.
  int lower = id * chunk;                      //   ⌈m/p⌉ ; assign 1 more to earlier threads
  int upper = min(lower + chunk, m);
  for (int row = lower ~ upper-1)
  { double accum = 0.0;
    for (int col = 0 ~ n-1)
      accum += A[row*n + col] * x[col];
    b[row] = accum;
  }
}
```

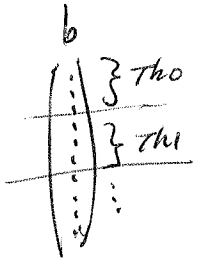


```
int main(---)
{ :
  double* A = new double[m*n];
  double* x = new double[n];
  double* b = new double[m];
  //initialize A and x here
  vector<thread> threads;
  for (int id = 0 ~ num_threads-1)
    threads.emplace_back(block_parallel_mult, id, A, x, b, m, n,
                         num_threads);
  for (auto& thread : threads)
    thread.join();          ——— wait for all threads terminate
  for (int i = 0 ~ m-1)
    cout << b[i] << endl;
  delete[] A; delete[] x; delete[] b;
  return 0;
}
```

— block distribution (DMV) — version 2 — using lambda function

```
:
void block_parallel_mult(double A[], double x[], double b[],
                         int m, int n, int num_threads)
{ auto block = [&](int id) -> void
  { int chunk = ceil((double) m / num_threads);
    int lower = id * chunk;
    int upper = min(lower+chunk, m);
    for(int row = lower ~ upper-1)
    { double accum = 0.0;
      for (int col=0 ~ n-1)
        accum += A[row*n + col] * x[col];
      b[row] = accum;
    }
  };
  vector<thread> threads;
  for(int id=0 ~ num_threads-1)
    threads.emplace_back(block, id);
  for(auto& thread : threads)
    thread.join();
}
```

(lambda func.)

```
int main(--)
{ :
  double* A = new double[m*n];
  double* x = ---
  double* b = ---
  // init A and x here.
  block_parallel_mult(A, x, b, m, n, num_threads);
  for(int i=0 ~ m-1)
    cout << b[i] << endl;
  delete[] A; delete[] x; delete[] b;
  return 0;
}
```

- Lambda function usage

ex)

- `auto add_one = [ ] (int& v) { return v+1; };`

error
- `int w = 1;`
  `auto add_w = [ ] (int& v) { return v+w; };` ← compile time error
  (w is declared out of λ-func.)

- `int w=1;`
  `auto add_w = [w] (int& v) { return v+w; };`
  capture w by value

- `int w=1;`
  `auto add_w = [&w] (int& v) { return v+w; };`
  capture w by reference

- `auto add_w = [&] (int& v) { return v+w; };`
  capture everything accessed in λ-func. by reference

- `auto add_w = [=] (int& v) { return v+w; };`
  capture everything accessed in λ-func. by value

cyclic distribution (PMV)

```
:
void cyclic_parallel_mult (double A[], double X[], double b[],
                           int m, int n, int num_threads)
{
  auto cyclic = [&] (int id) -> void
  { for (row = id ~ m-1, row += num_threads)
    { accum = 0.0;
      for (col = 0 ~ n-1)
        accum += A[row*n + col] * X[col];
      b[row] = accum;
    };
  vector<thread> threads;
  for (id = 0 ~ num_threads - 1)
    threads.emplace_back (cyclic, id);
  for (auto& thread : threads)
    thread.join();
}
```

λ-func

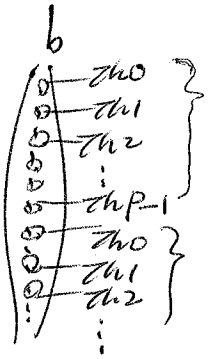

fine-grained cyclic

may cause → false sharing issue

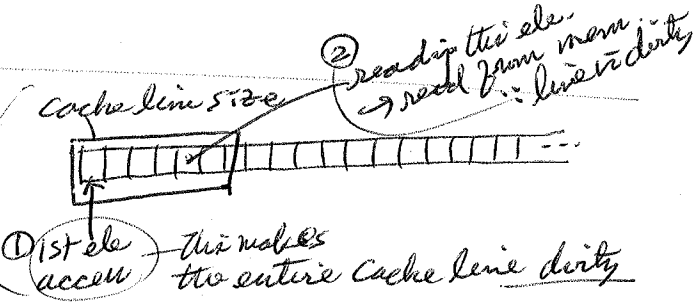

→ Sol: block-cyclic distribution

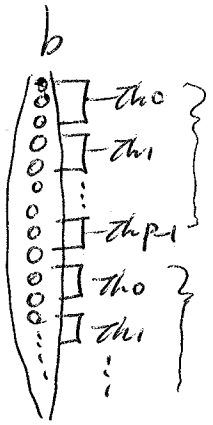

in fact,

- block distr — block size = $\left(\frac{m}{p}\right)$
- cyclic distr — block size = 1

extreme cases of block-cyclic distr.

— block size should be same as cache line size.

false sharing

ex)
```
auto cyclic = [&b] (int id) -> void
{ for (int row=id; row<m; row+= num_threads)
  { b[row]=0;
    for (col=0~n-1)
      b[row] += A[row*n+col] * X[col];
  }
};
```

→ result is correct, but memory access overheads (due to false sharing).

ex) b[0] is updated by Th0 on core0 ; ← cache line in core0 is dirty
b[1] is read by Th1 on core1 ; ← core1 accesses mem for same cache line. (or updated cache)

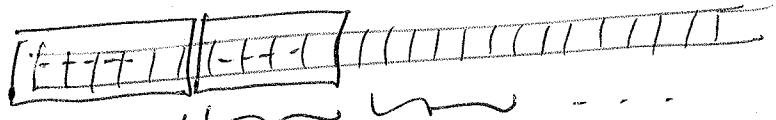

(block size = cache line size) accessed by Th0 — accessed by Th1

no conflict

ex) Intel
cache line size 64 bytes
⇒ block size
good { c=16 for int
c=8 for 64bit data

ex) Thi — local cache update and then, read same line.
→ no problem — cache up to date.

— block size
(too) small
→ false sharing
vs.
→ load balanced
big
→ cache friendly access
→ load imbalance

block-cyclic distribution (DMV)

```
:
void block_cyclic_parallel_mult(double A[], ... X[], ... b[],
                                int m, int n, int num_threads, int chunk_size)
{ auto block_cyclic = [&](int id) -> void
  { int offset = id * chunk_size;   // chunk_size = 64/8 = 8
    int stride = num_threads * chunk_size;   // next jump
    for (int lower = offset; lower < m; lower += stride)
    { int upper = min(lower + chunk_size, m);
      for (int row = lower; row < upper; row++)
      { double accum = 0.0;
        for (int col = 0 ~ n-1)
          accum += A[row*n + col] * X[col];
        b[row] = accum;
      }
    }
  };// auto
  vector<thread> threads;
  for (int id = 0 ~ num_threads - 1)
    threads.emplace_back(block_cyclic, id);   // call λfunc.
  for (auto& thread : threads)
    thread.join();
}

int main(---)
{ :
  block_cyclic_parallel_mult(A, x, b, m, n, num_threads, 64/sizeof(double));   // = 8
  for (i = 0 ~ m-1)
    cout << b[i] << endl;
:
}
```

(64 = cache line size; sizeof(double) = double type size)